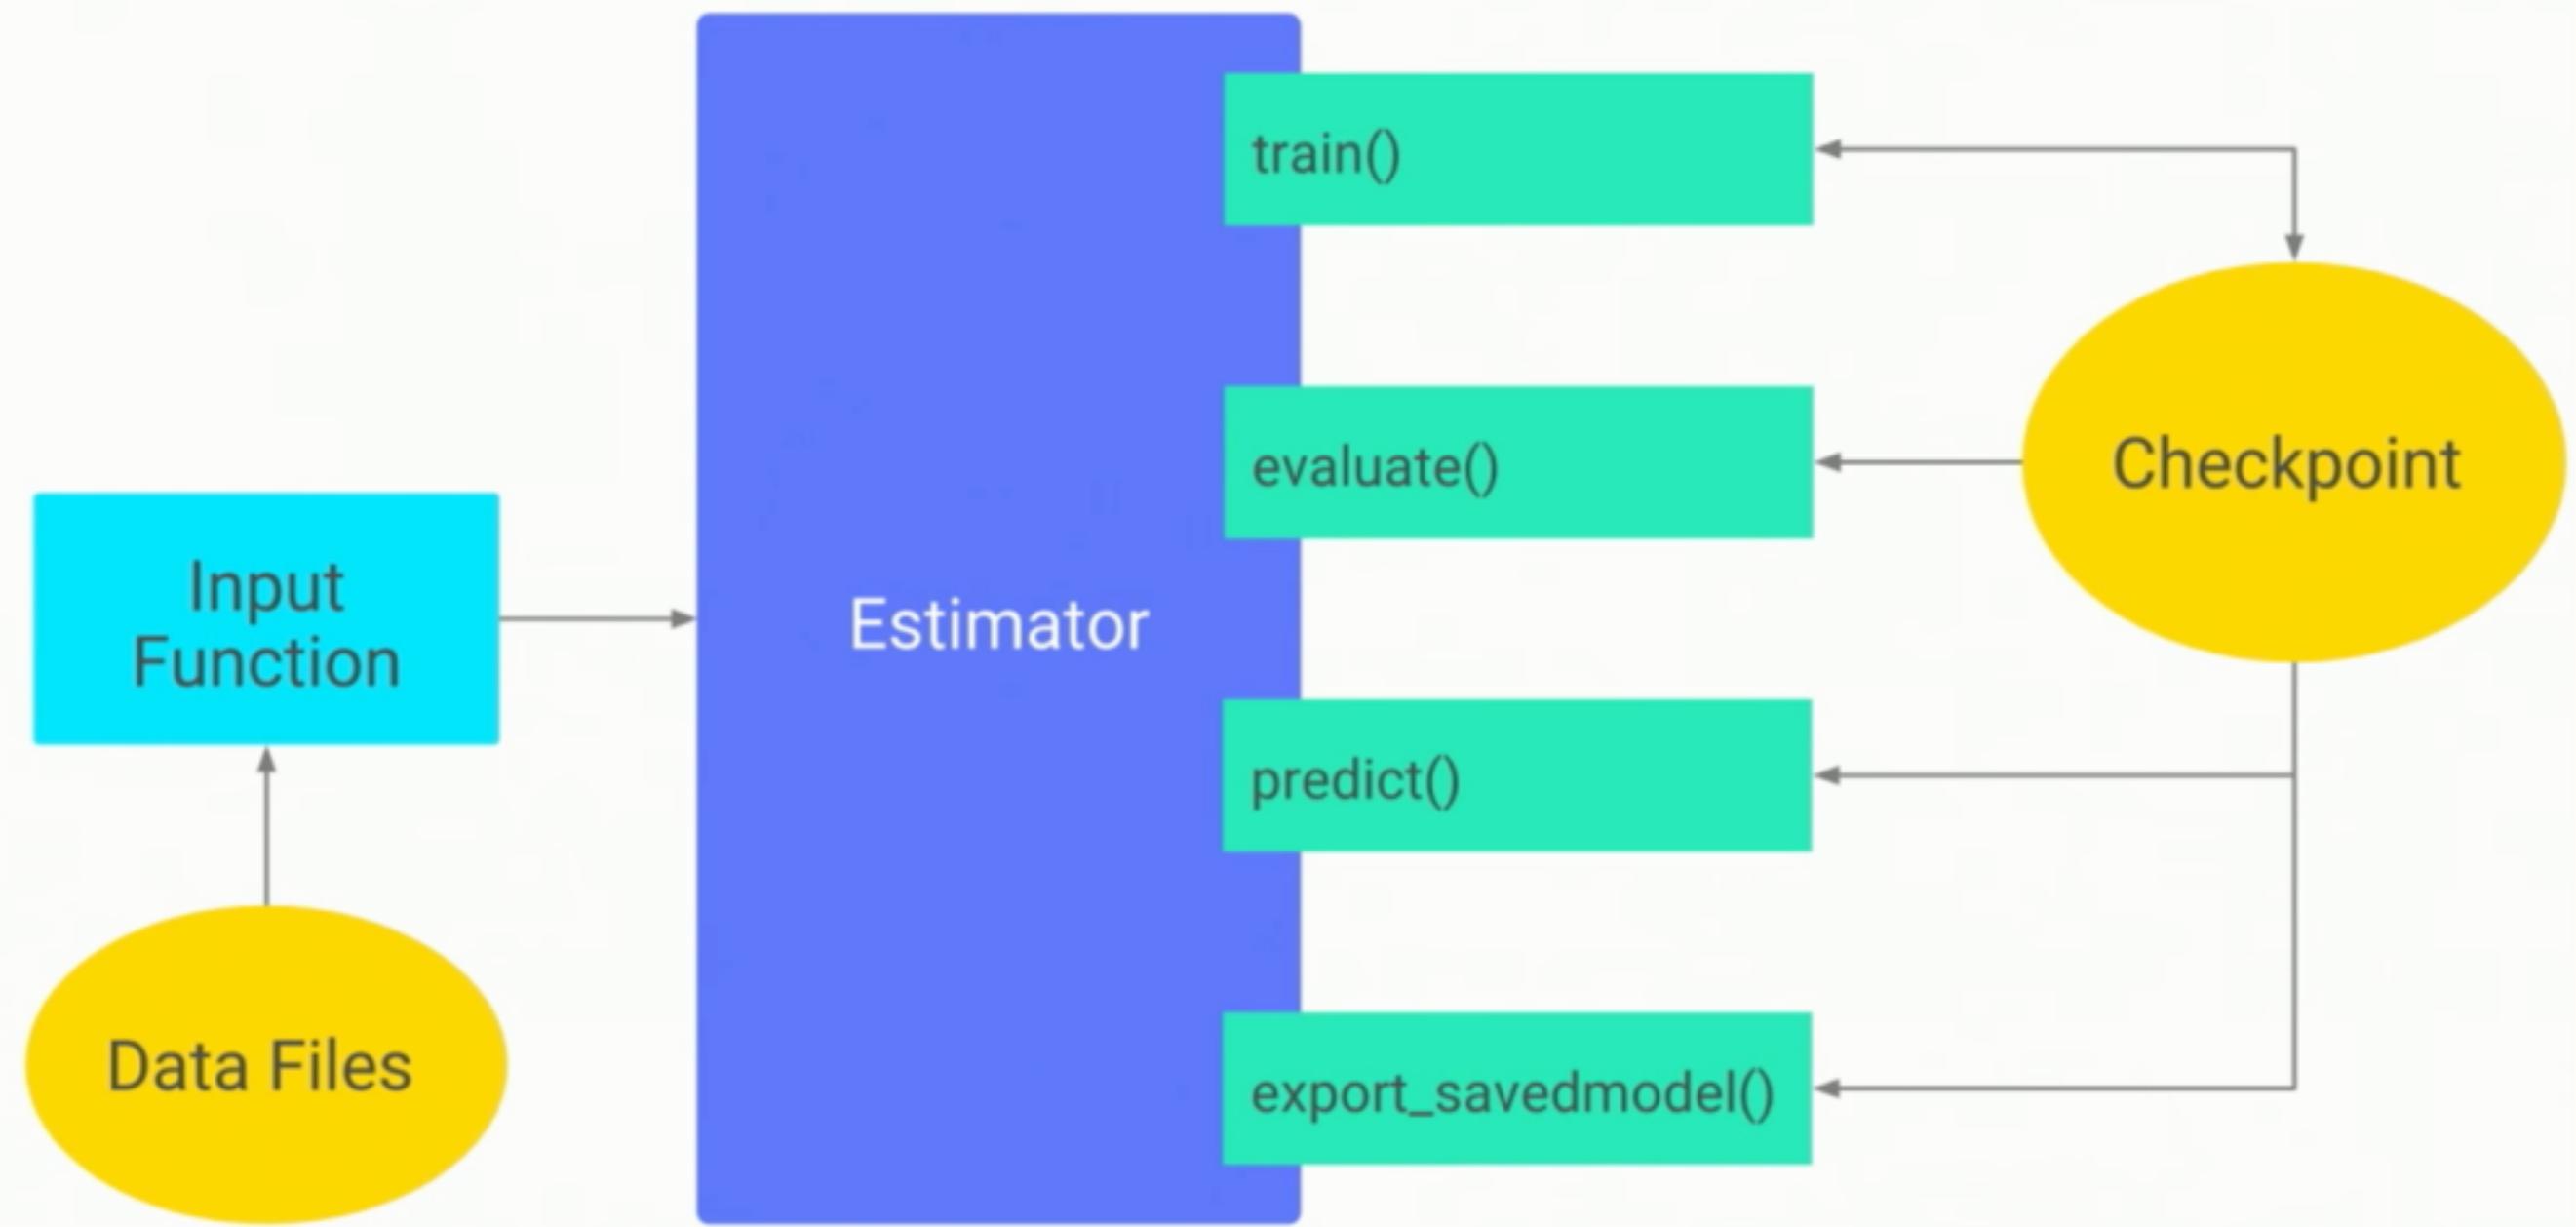
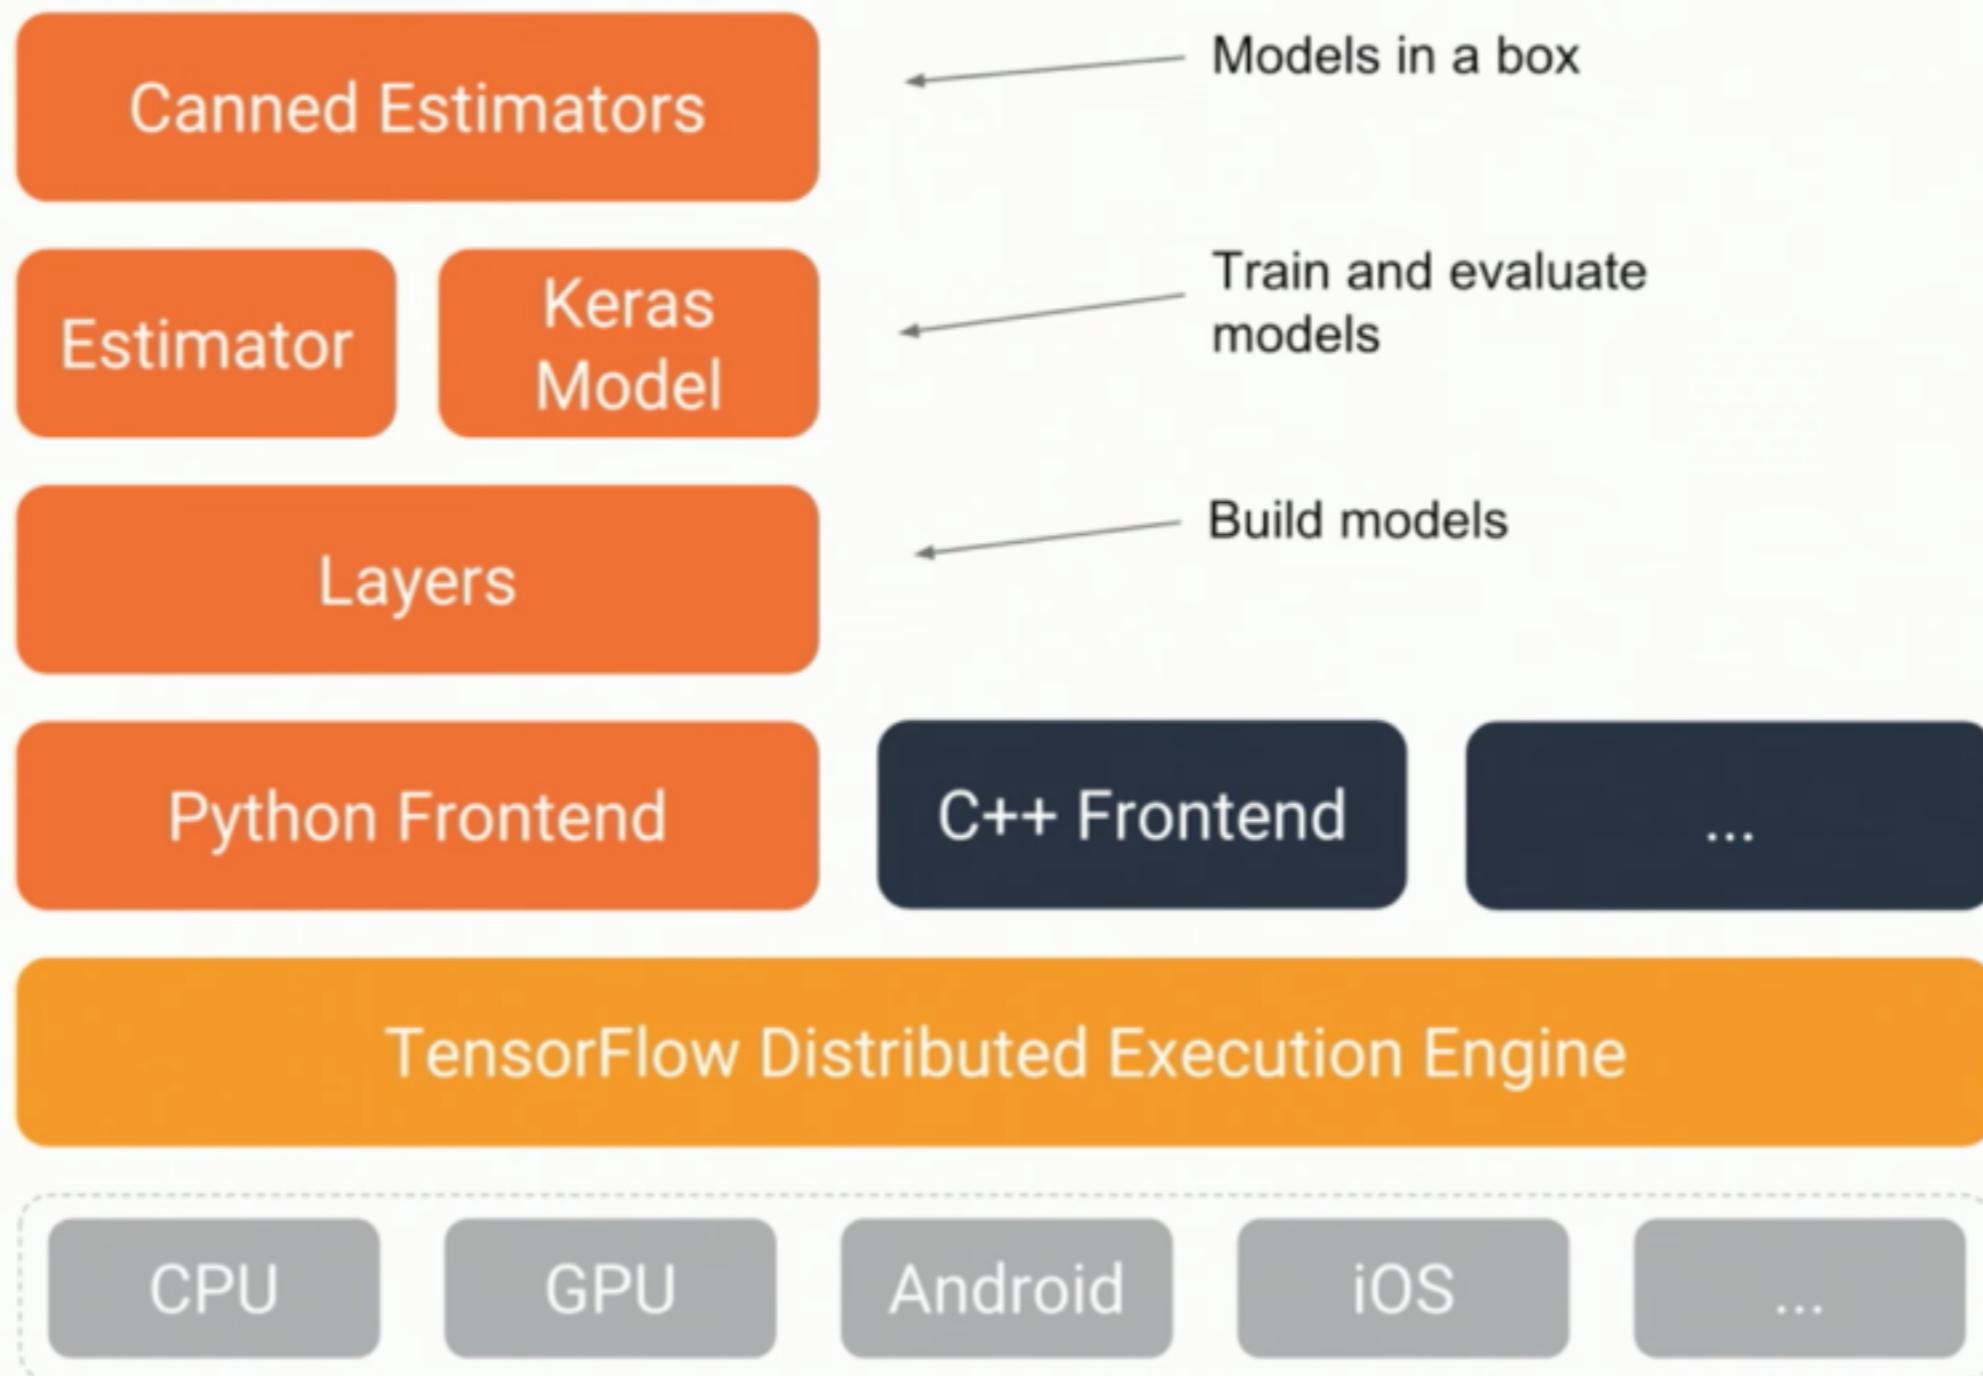


Effective TensorFlow 非专家入门指南

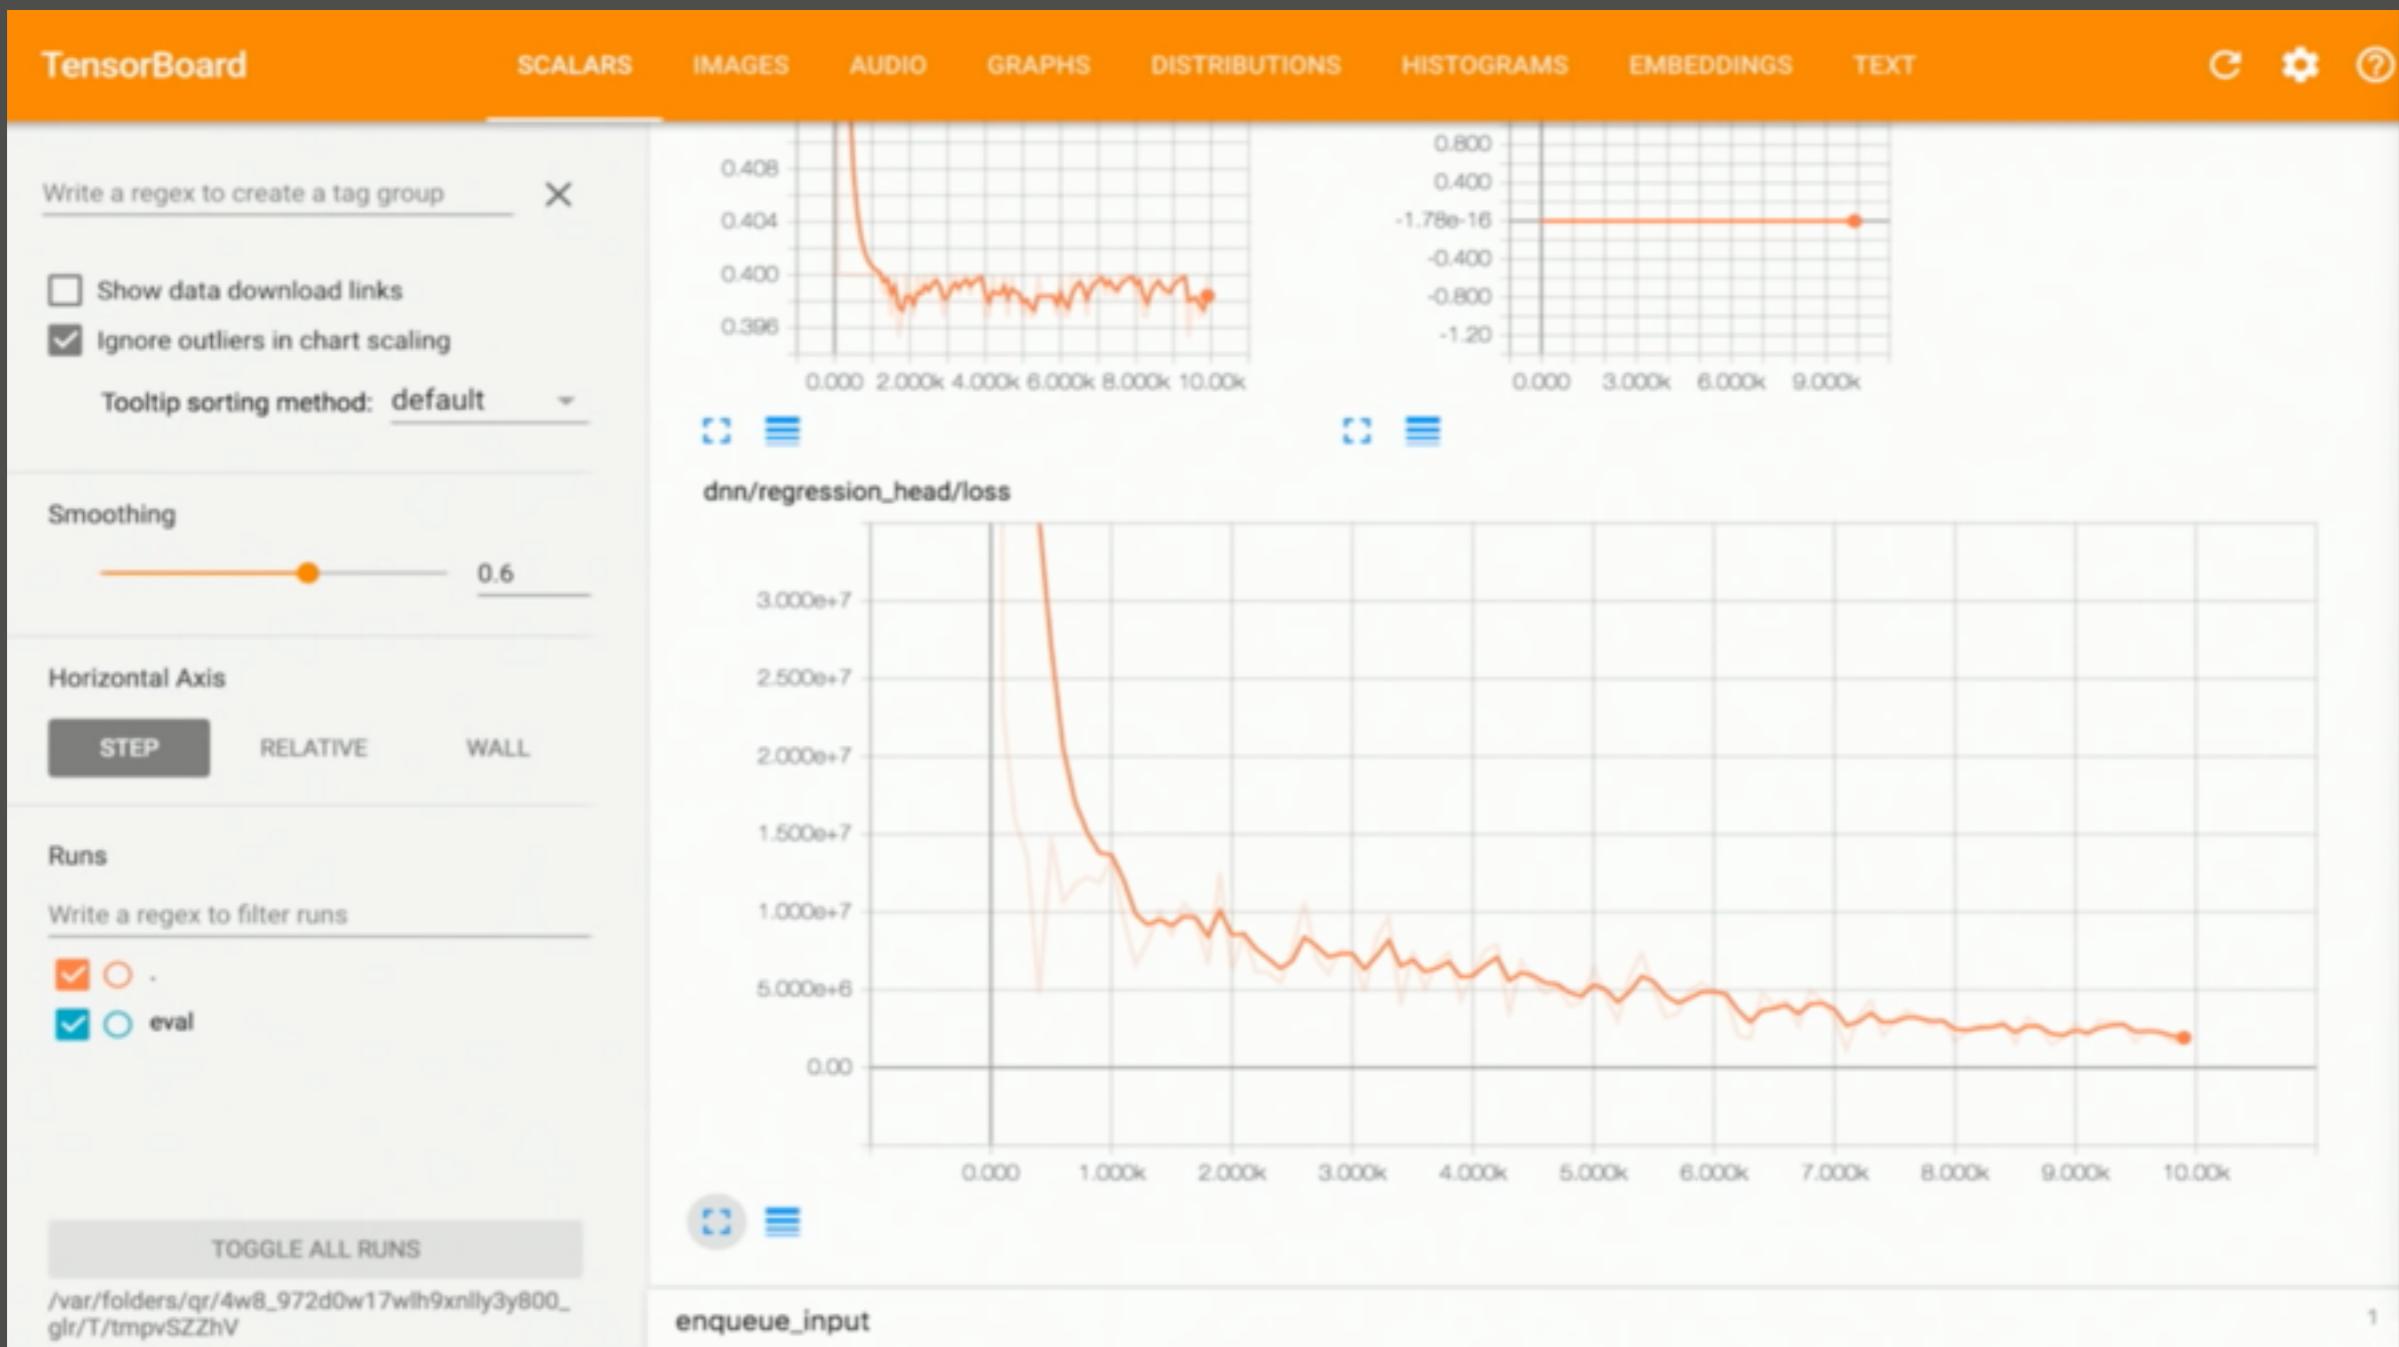
廖桔秋



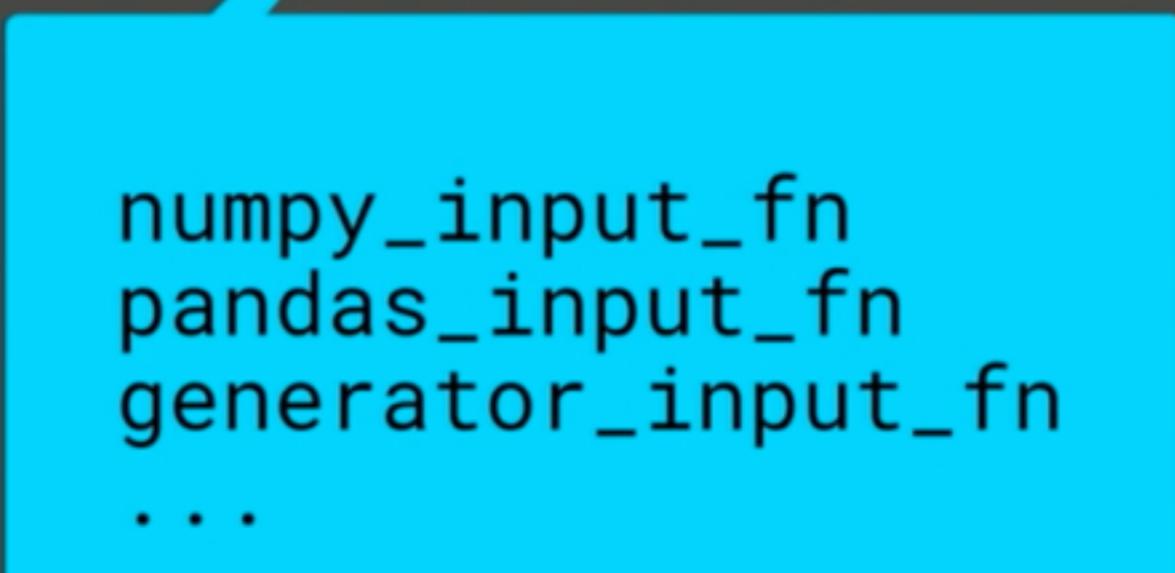



```
train_input_fn = pandas_input_fn(x=input_data, y=input_label,  
                                batch_size=64, shuffle=True,  
                                num_epochs=None)  
  
regressor.train(train_input_fn, steps=10000)
```

INFO:tensorflow:loss = 6.07866e+07, step = 102 (0.884 sec)
INFO:tensorflow:global_step/sec: 131.905
INFO:tensorflow:loss = 7.26533e+07, step = 202 (0.758 sec)
INFO:tensorflow:global_step/sec: 125.879
INFO:tensorflow:loss = 3.64895e+07, step = 302 (0.794 sec)
INFO:tensorflow:global_step/sec: 131.558
INFO:tensorflow:loss = 4.84076e+07, step = 402 (0.760 sec)
INFO:tensorflow:global_step/sec: 124.394
INFO:tensorflow:loss = 4.316e+07, step = 502 (0.804 sec)
INFO:tensorflow:global_step/sec: 126.16
INFO:tensorflow:loss = 2.32498e+07, step = 602 (0.793 sec)
INFO:tensorflow:global_step/sec: 122.751
INFO:tensorflow:loss = 4.53009e+07, step = 702 (0.817 sec)
INFO:tensorflow:global_step/sec: 118.297
INFO:tensorflow:loss = 4.56962e+07, step = 802 (0.843 sec)
INFO:tensorflow:global_step/sec: 126.641
INFO:tensorflow:loss = 3.3779e+07, step = 902 (0.790 sec)
INFO:tensorflow:global_step/sec: 109.258
INFO:tensorflow:loss = 5.32131e+07, step = 1002 (0.915 sec)
INFO:tensorflow:global_step/sec: 135.997
INFO:tensorflow:loss = 3.2445e+07, step = 1102 (0.735 sec)
INFO:tensorflow:global_step/sec: 123.69
INFO:tensorflow:loss = 2.43225e+07, step = 1202 (0.811 sec)
INFO:tensorflow:global_step/sec: 136.363
INFO:tensorflow:loss = 2.86487e+07, step = 1302 (0.730 sec)



```
train_input_fn = pandas_input_fn(x=input_data, y=input_label,  
                                 batch_size=64, shuffle=True,  
                                 num_epochs=None)
```

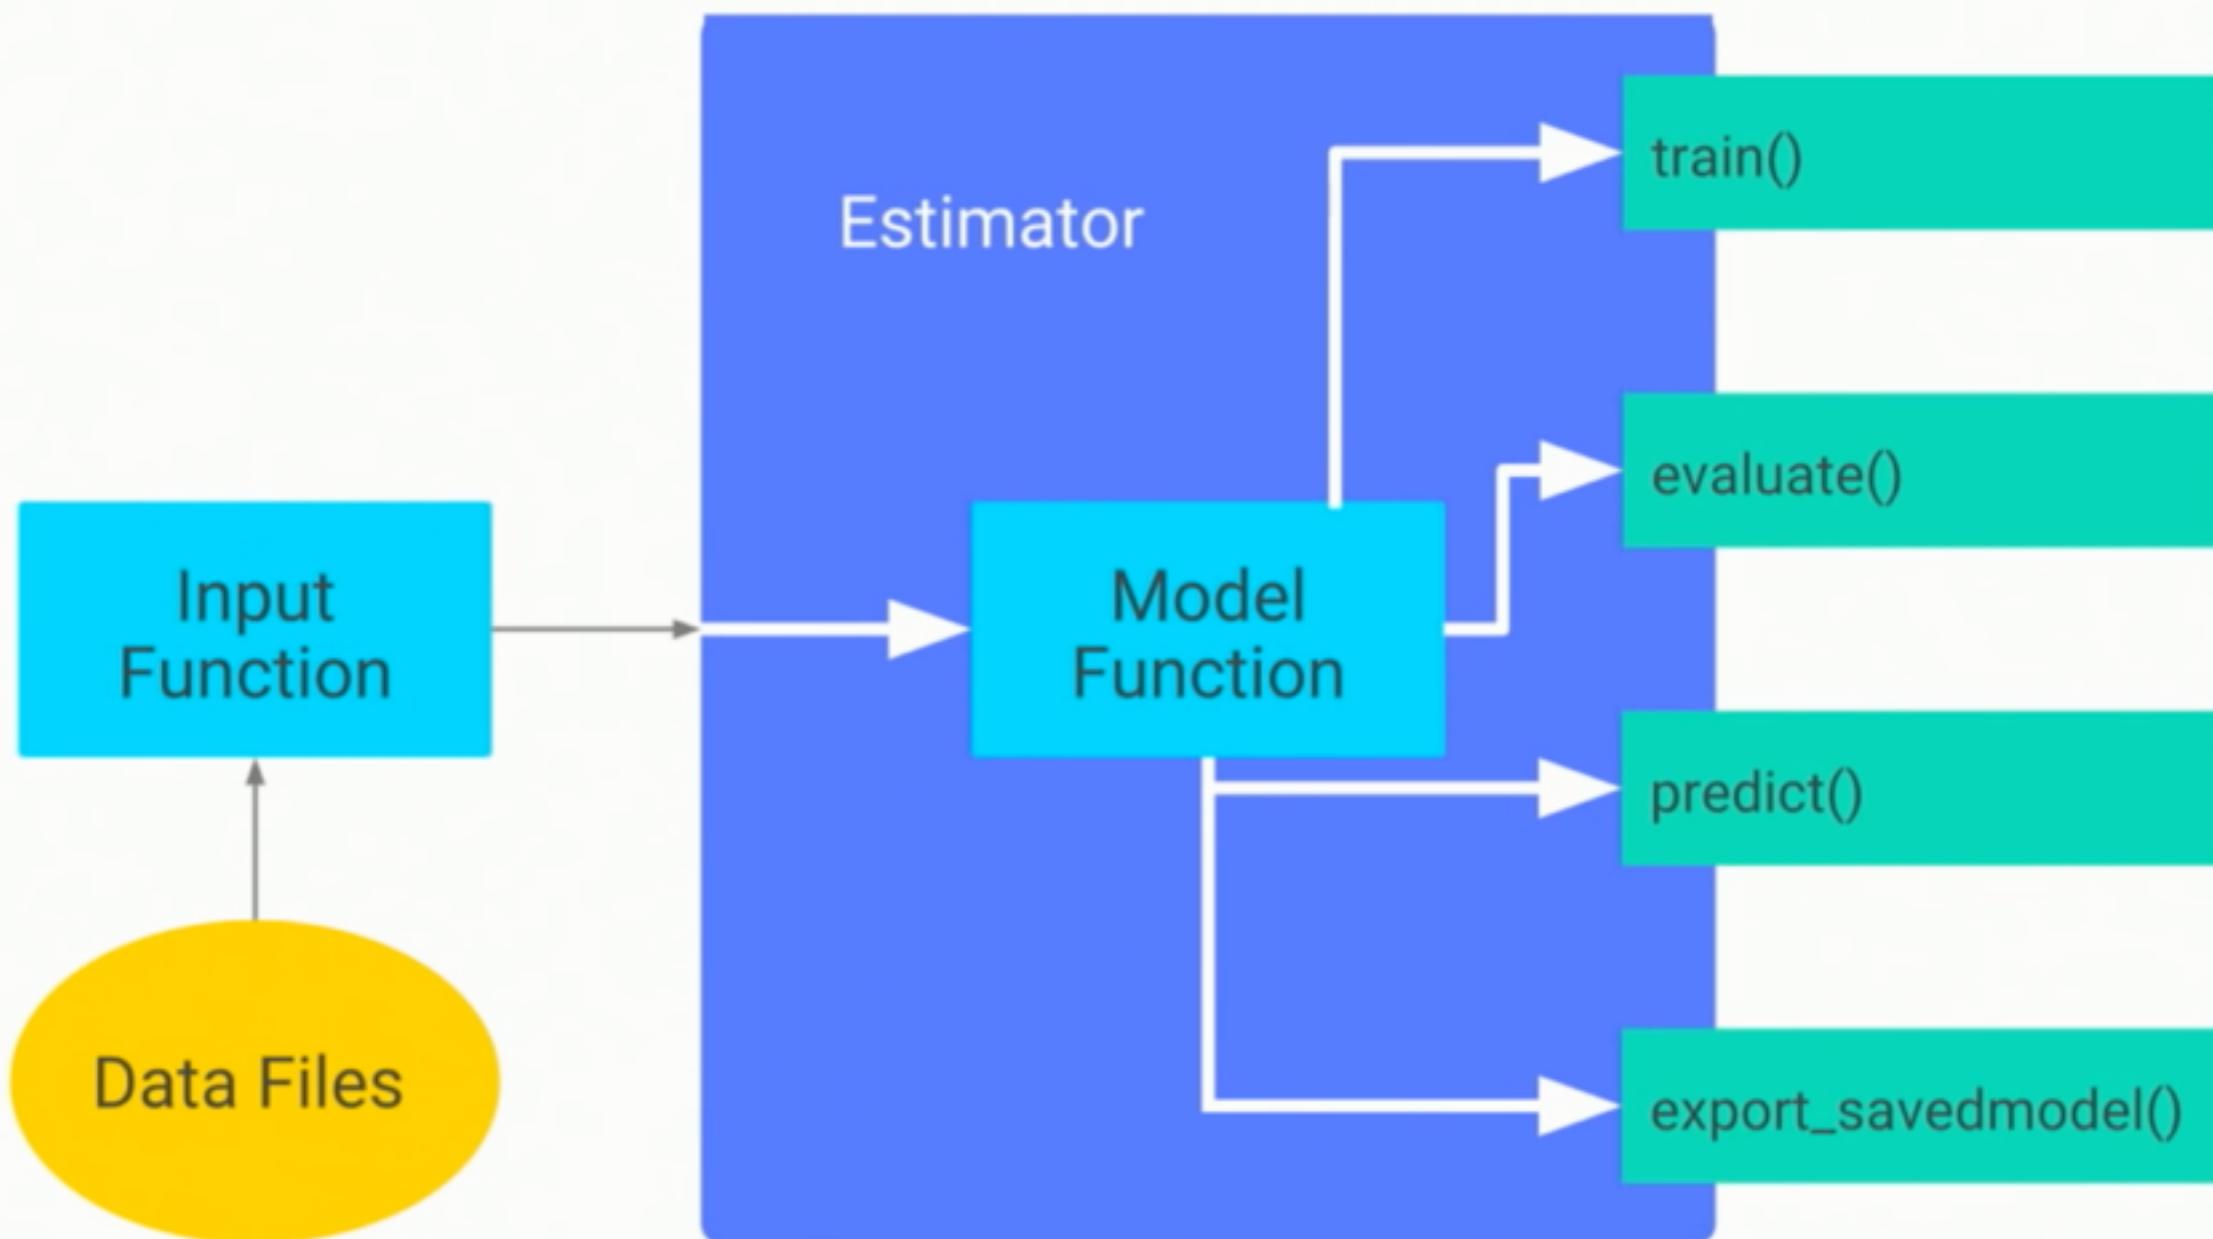


```
numpy_input_fn  
pandas_input_fn  
generator_input_fn  
...
```

```
regressor = LinearRegressor(feature_columns=[make, horsepower,  
                                             cylinders, ...])
```

LinearClassifier
DNNRegressor
DNNClassifier
TensorForest
KMeans
...

```
cylinders = categorical_column_with_vocabulary_list(  
    'num-of-cylinders', ['two', 'three', 'four', ...])  
...  
  
regressor = DNNRegressor(feature_columns=[embedding_column(make, 10),  
                                         horsepower,  
                                         indicator_column(cylinders, 3),  
                                         ...],  
                         hidden_units=[50, 30, 10])
```



```
def experiment_fn(run_config, hparams):
    # Make Estimator
    regressor = DNNRegressor(..., config=run_config,
                            hidden_units=hparams['units'])

    # Collect information for training
    return Experiment(estimator=regressor,
                      train_input_fn=pandas_input_fn(...),
                      eval_input_fn=pandas_input_fn(...))

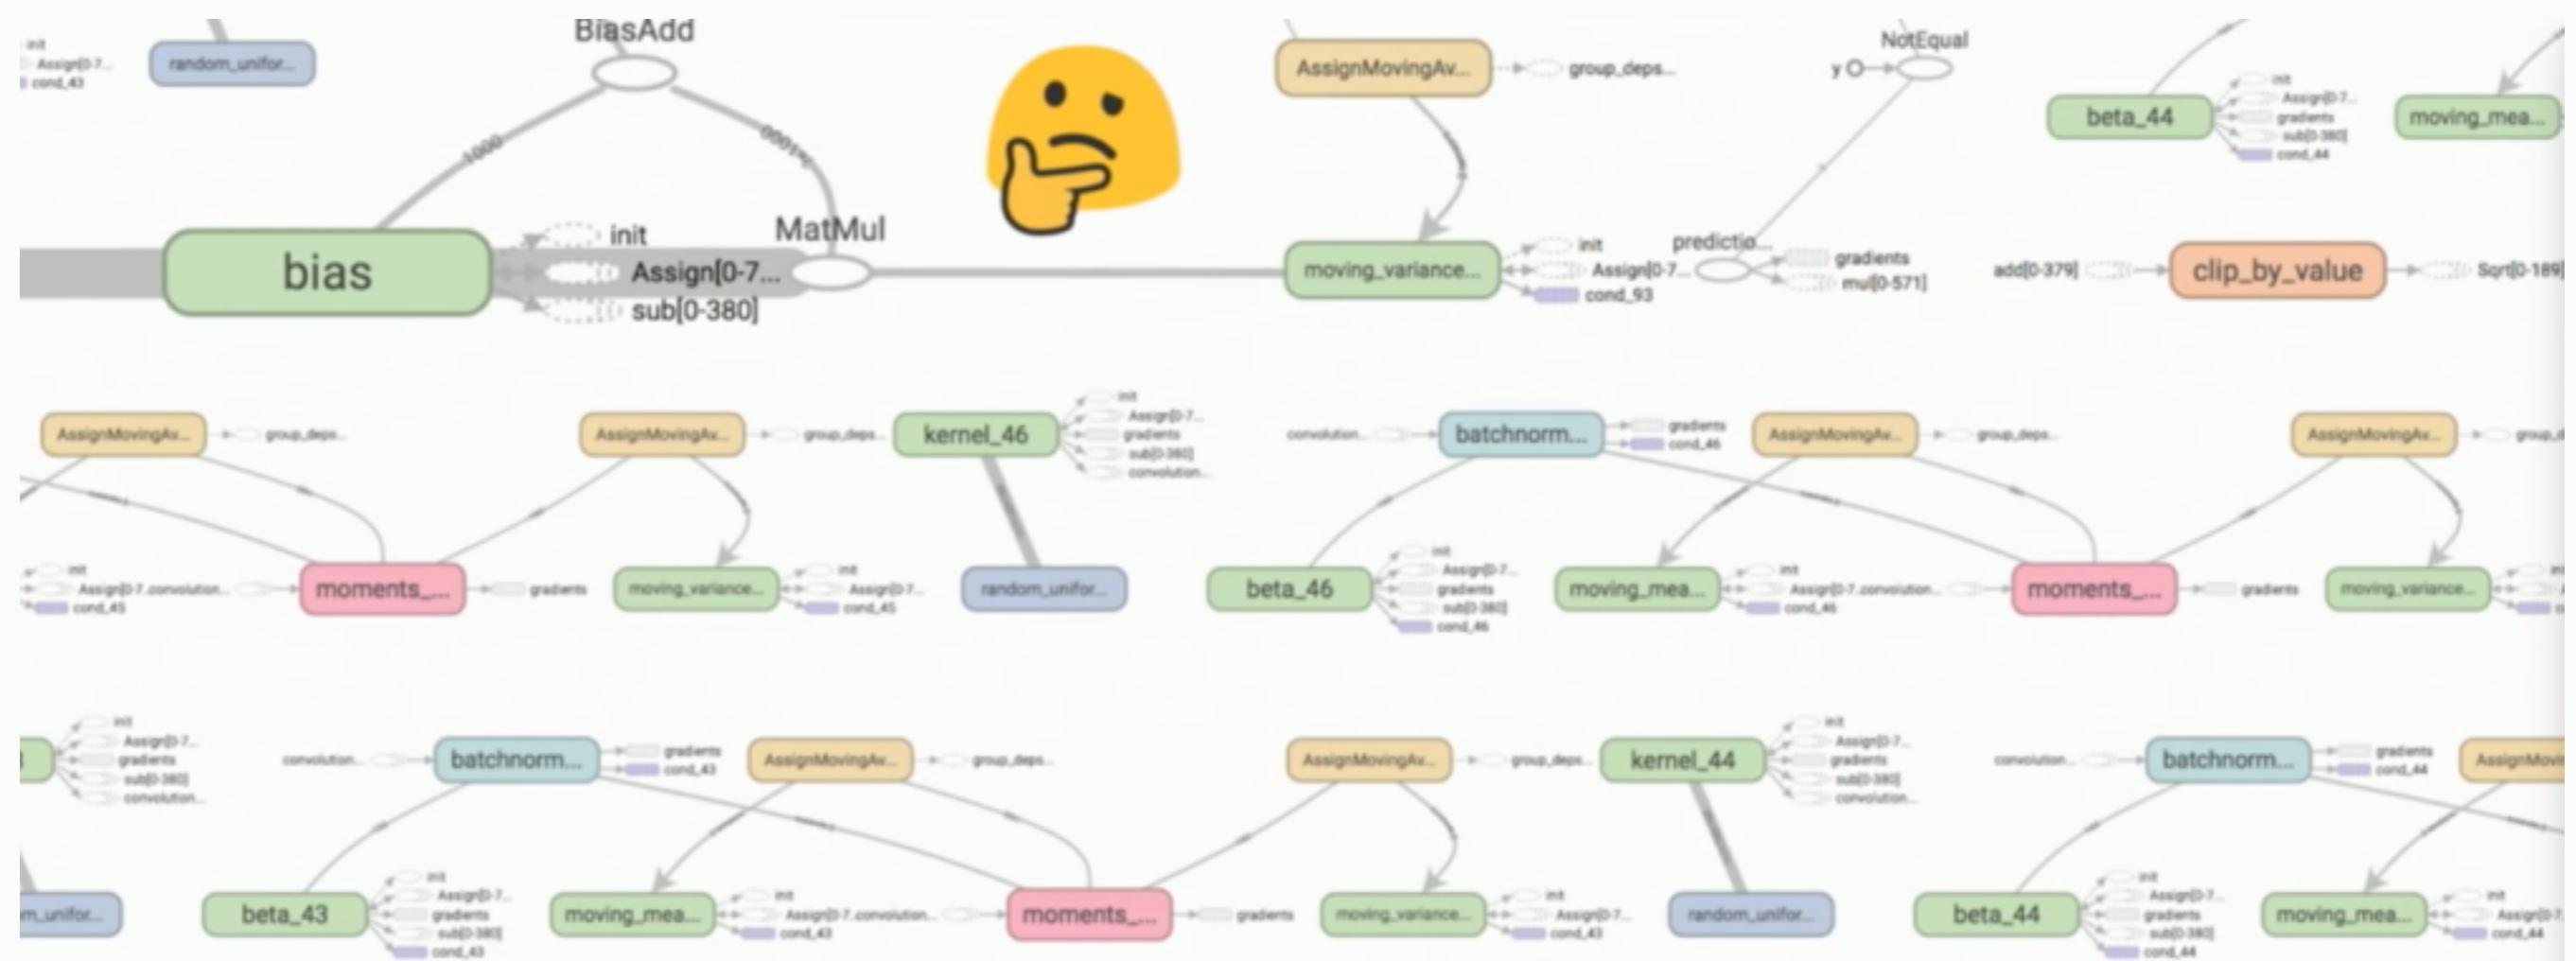
if __name__ == '__main__':
    learn_runner.run(experiment_fn,
                     run_config=RunConfig(model_dir="/tmp/output_dir"))
```

```
cluster_spec = {
    'ps': ['host1:2222', 'host2:2222', ...],
    'worker': ['host3:2222', 'host4:2222', ...],
}

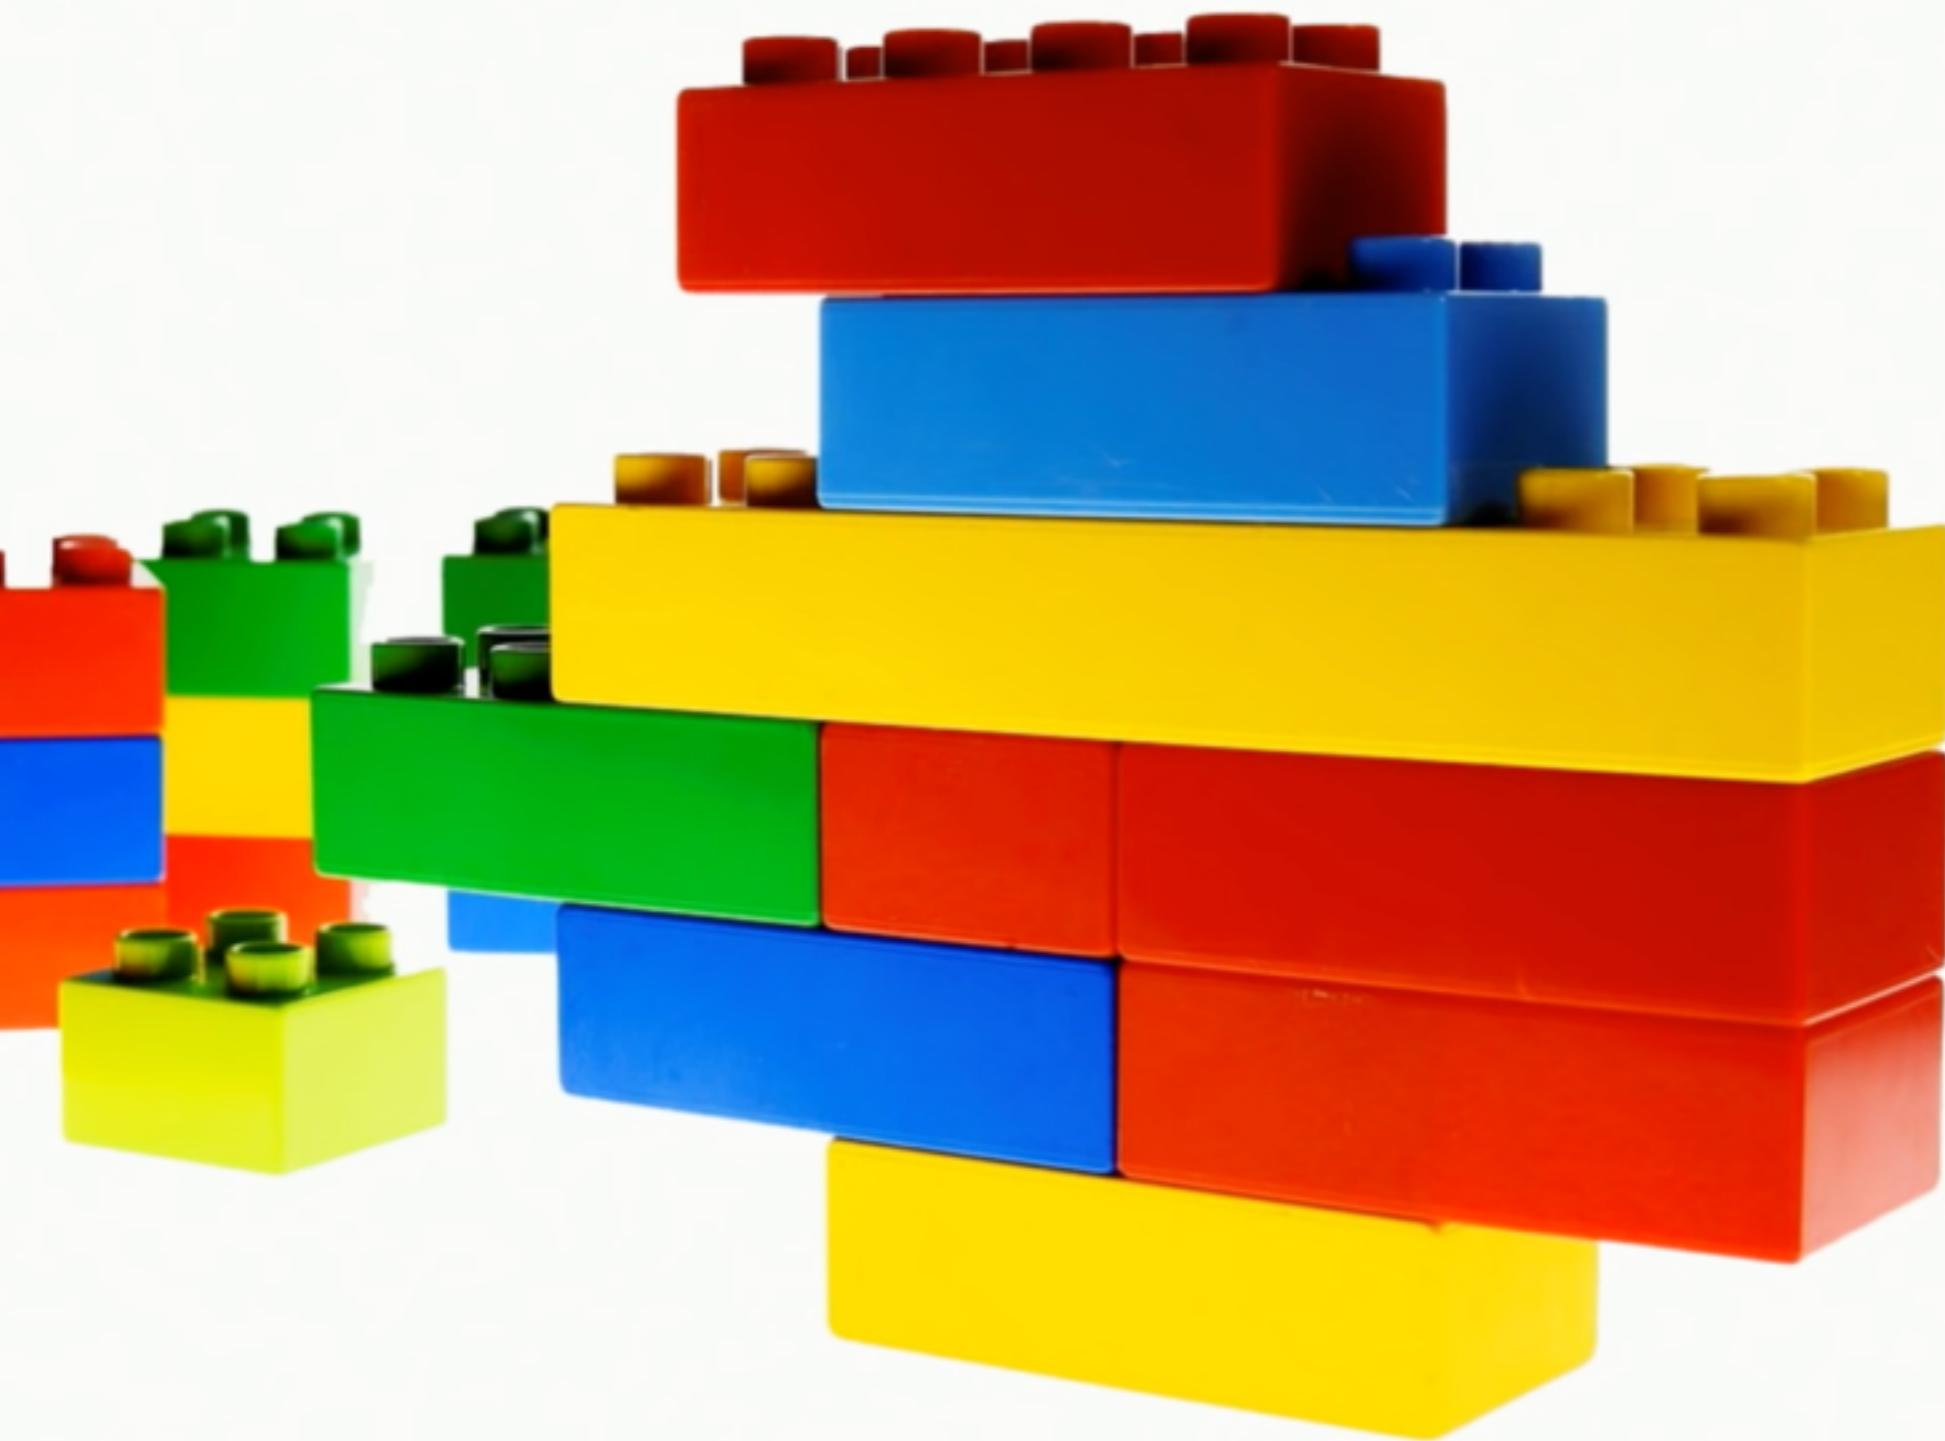
os.environ['TF_CONFIG'] = json.dumps({
    'cluster': cluster_spec,
    'task': {'type': 'worker', 'index': 1}
})
```

<https://github.com/tensorflow/ecosystem>

- Complete ML models quickly
- Distributed training



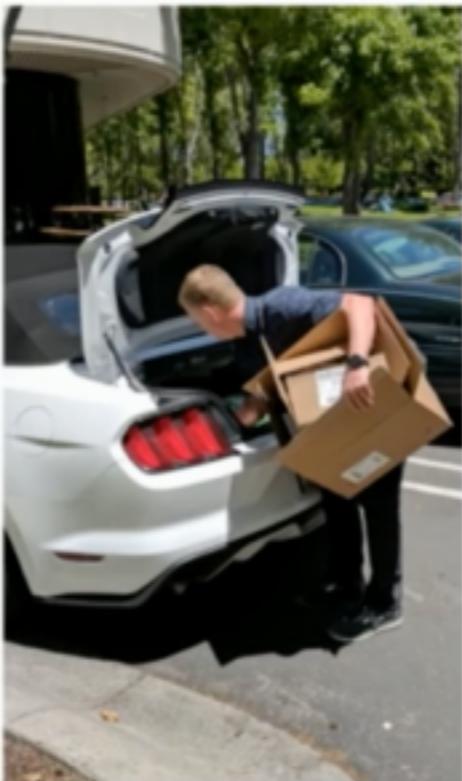
Building TensorFlow models from scratch



- Keras: Deep learning for everyone
- A high-level API specification
- Focus on User Experience

- Keras API as `tf.keras`
- For TensorFlow Users: a simplified workflow
- For Keras users: access to more powerful features: distributed training, hyper parameter optimization, training on Cloud ML

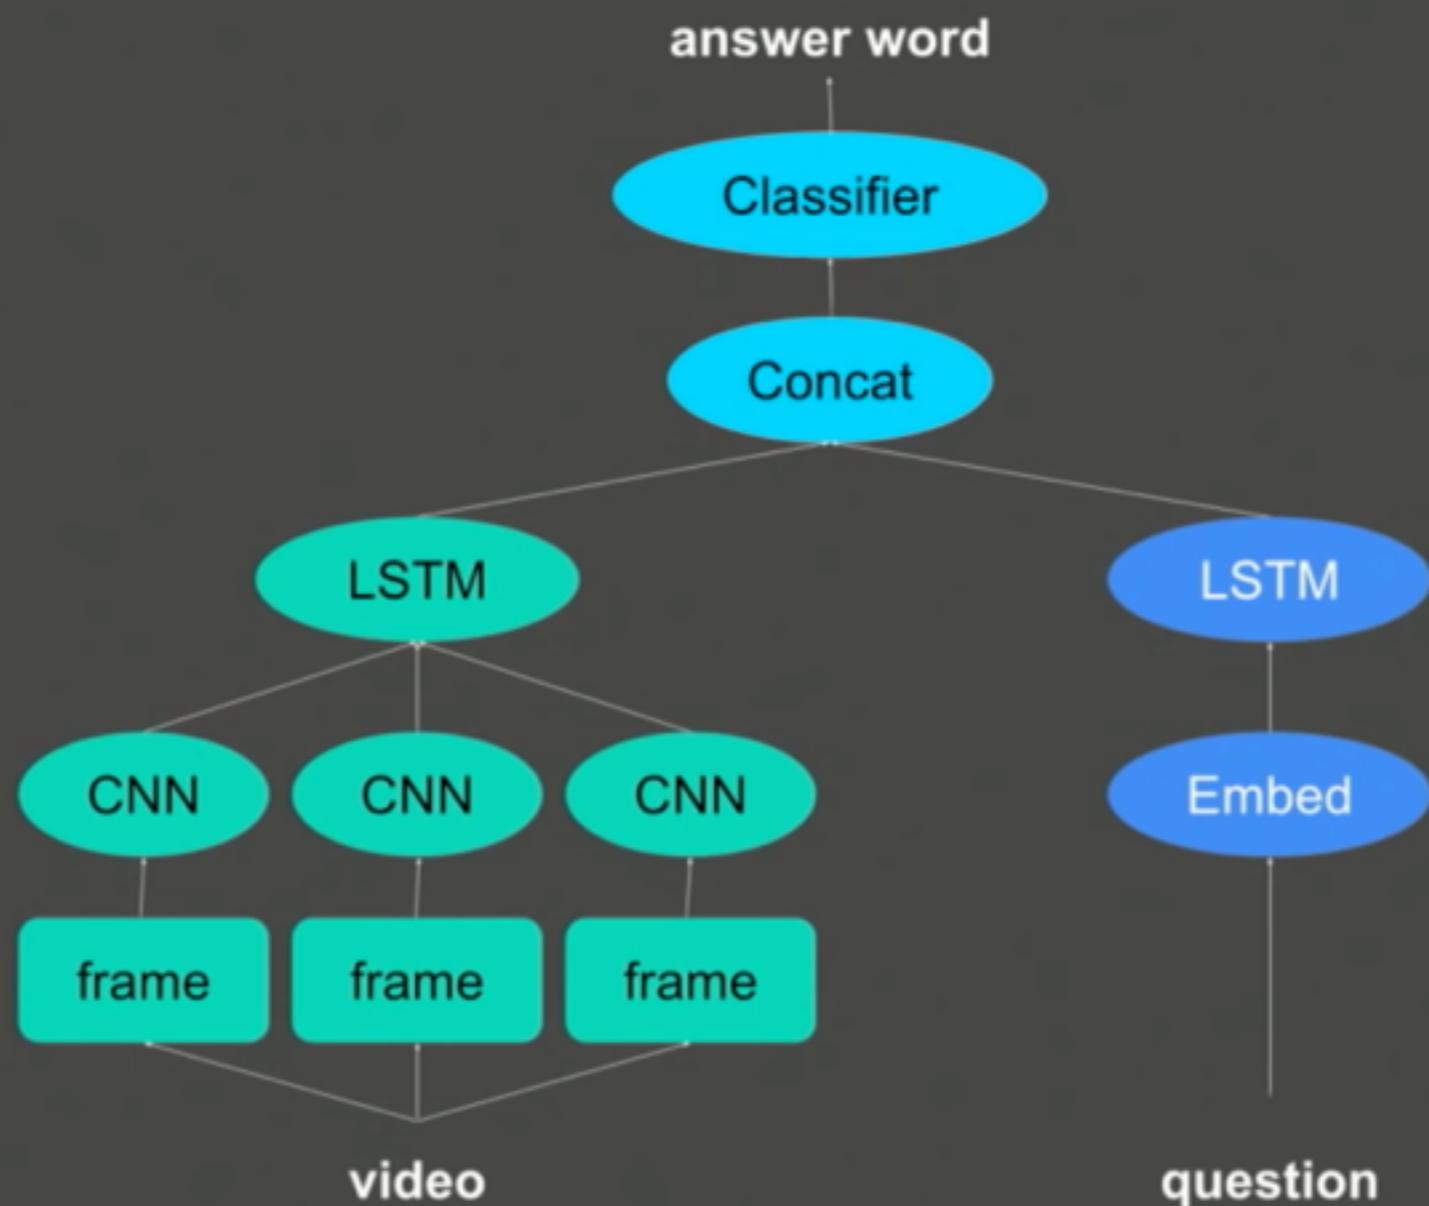
Toy video-QA problem

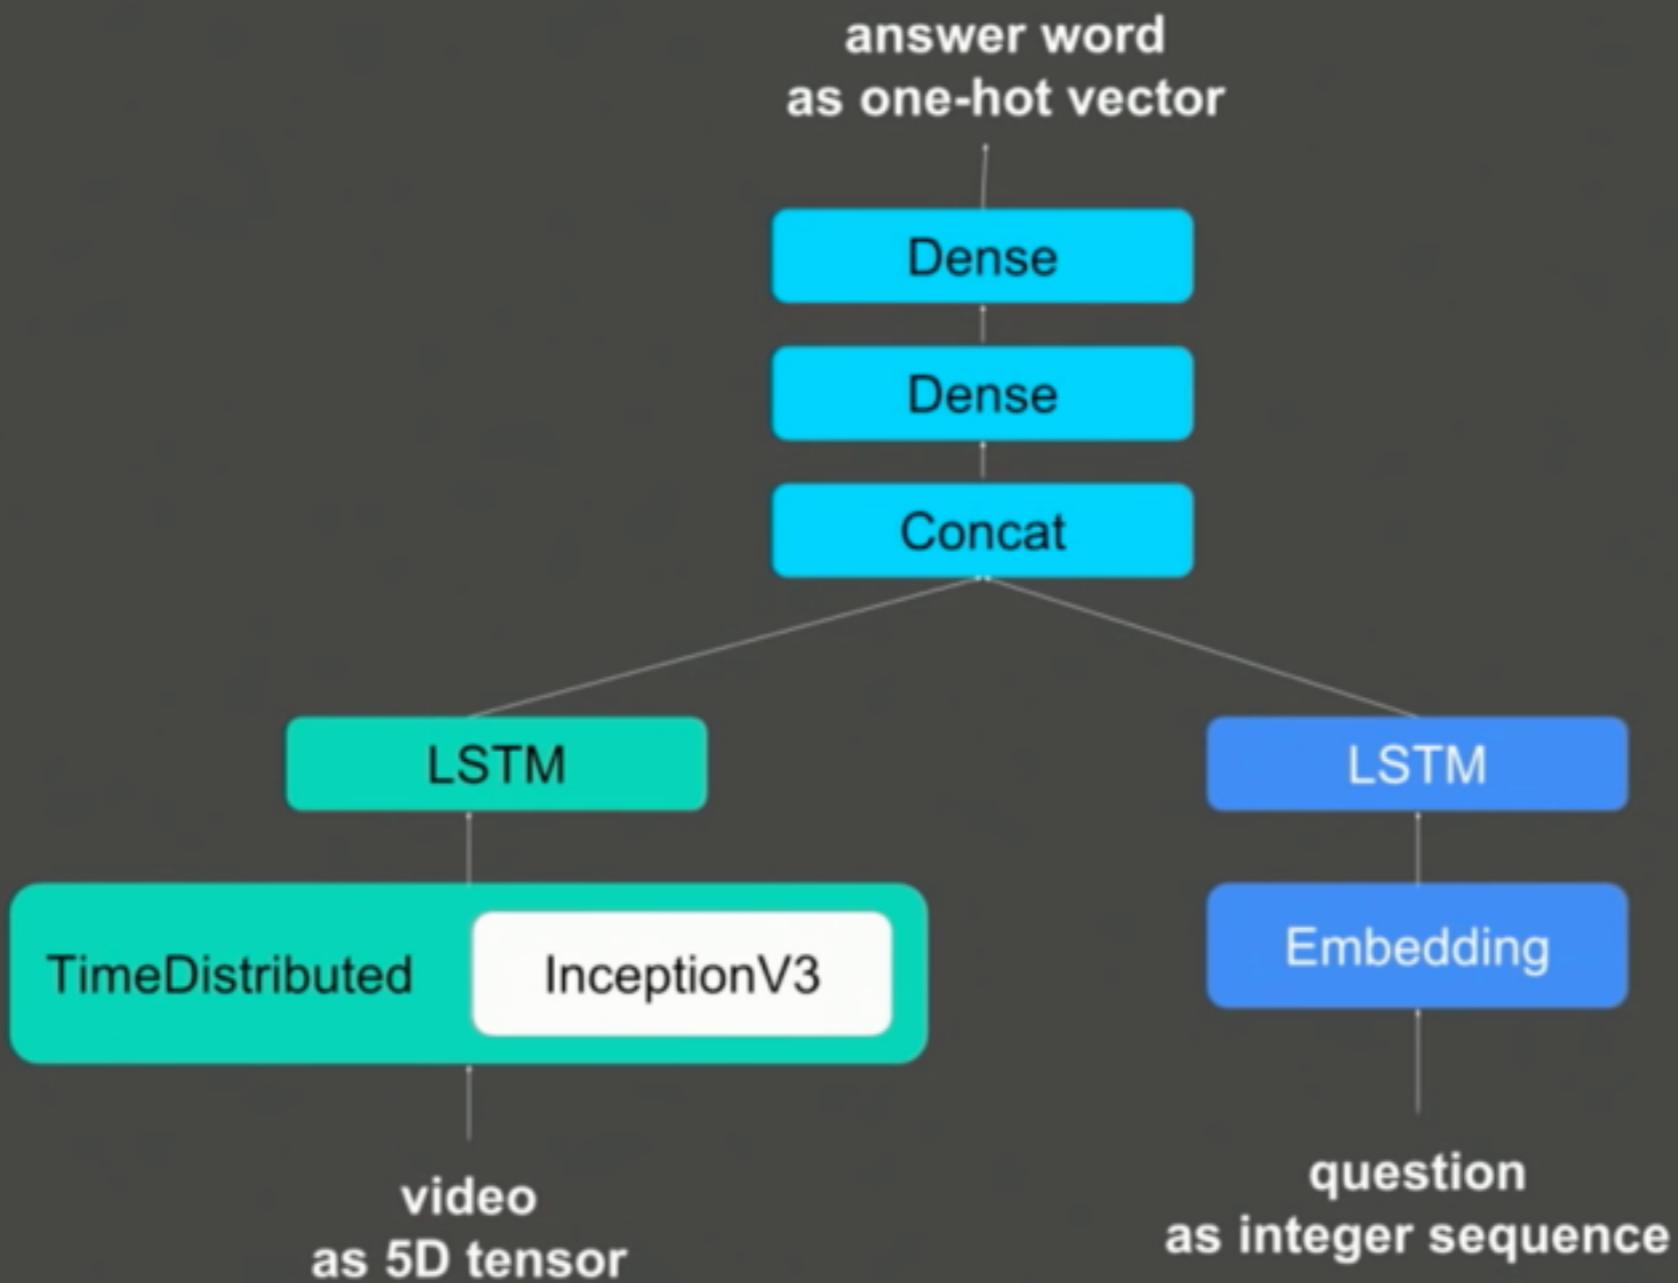


- > "**What is the man doing?**"
> **packing**

- > "**What color is his shirt?**"
> **blue**

A previously very hard problem,
made accessible to anyone with basic Python scripting
abilities.





Turning frames into a vector,
with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')
cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning a sequence of words into a vector

```
question = keras.Input(shape=(None,), dtype='int32', name='question')
embedded_words = layers.Embedding(input_voc_size, 256)(question)
question_vector = layers.LSTM(128)(embedded_words)
```

Best practice is the default config.

Predicting an answer word

```
x = layers.concatenate([video_vector, question_vector])
x = layers.Dense(128, activation=tf.nn.relu)(x)
predictions = layers.Dense(output_voc_size, name='predictions')(x)
```

Setting up the training configuration

```
model = keras.models.Model([video, question], predictions)
model.compile(optimizer=tf.AdamOptimizer(),
              loss=tf.softmax_crossentropy_with_logits)
```

Leveraging Experiment for distributed training

```
def experiment_fn(config, params):
    model = ...
    estimator = model.get_estimator(config=...)
    return Experiment(estimator,
                      train_input_fn=pandas_input_fn(...),
                      eval_input_fn=pandas_input_fn(...))

if __name__ == '__main__':
    tf.contrib.train.run_experiment(experiment_fn)
```

- Eras: a high-level API with focus on UX.
- Now part of TensorFlow
- To use with Estimators and Experiments for distributed training

A big step towards
making deep learning accessible to everyone.